



(a) Latent action models

Which reconstruction target yields transferable actions?



(b) NWM architecture and pretraining

Alternative action interfaces · same backbone architecture

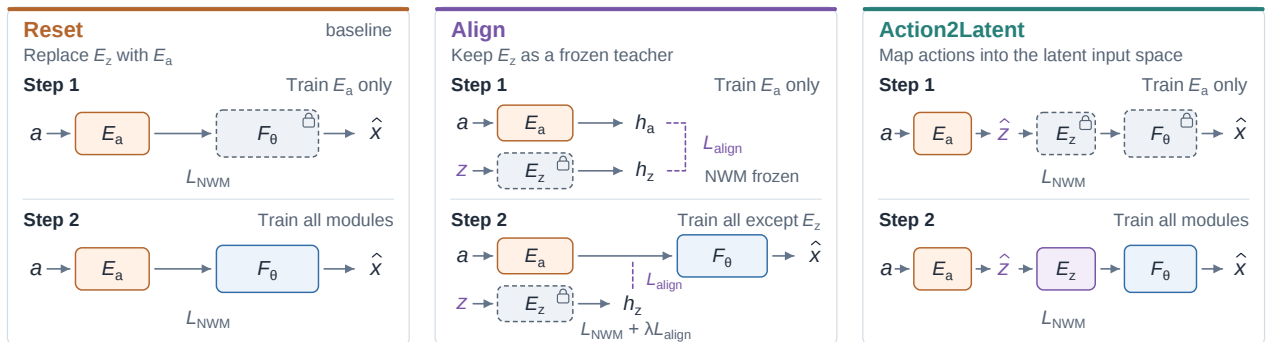


The same context and time / action conditions guide each denoising step. u : visual latent; z : latent action.

(c) Two-stage fine-tuning

From nwm-latentpt · supervised on Base data

Trainable Frozen



Training-only teacher in Align; Action2Latent retains E_z at inference. Visual / diffusion inputs are omitted in (c).

EVALUATION Visual prediction · Navigation planning · Unseen-domain generalization