

(b) NWM architecture and pretraining

Alternative action interfaces · same backbone architecture

Base data · GT navigation actions

RECON · SCAND · TartanDrive · HuRoN

NavAnywhere v1

13 datasets · 322 h

NavAnywhere v2 · scale extension

v1 + Ego4D + GO · 1200+ h

nwm-real

GT actions · no pretraining



Train on Base from scratch

nwm-timept

Pretrain on NavAnywhere v1

No action branch

Base: Reset fine-tuning

nwm-geopt

Pretrain on NavAnywhere v1



Base: Reset fine-tuning

nwm-latentpt

Pretrain on NavAnywhere v1




Base: strategies in (c)

Dashed arrows: offline video labeling




NWM F_θ : visual latent diffusion



GENERATION

at inference



The same context and time / action conditions guide each denoising step. u: visual latent; z: latent action.